


# Deep Learning Software

## Tensorflow/PyTorch

(Based on stanford cs231n slides)

# A zoo of frameworks!




We'll focus on these

# Choose a deep learning software

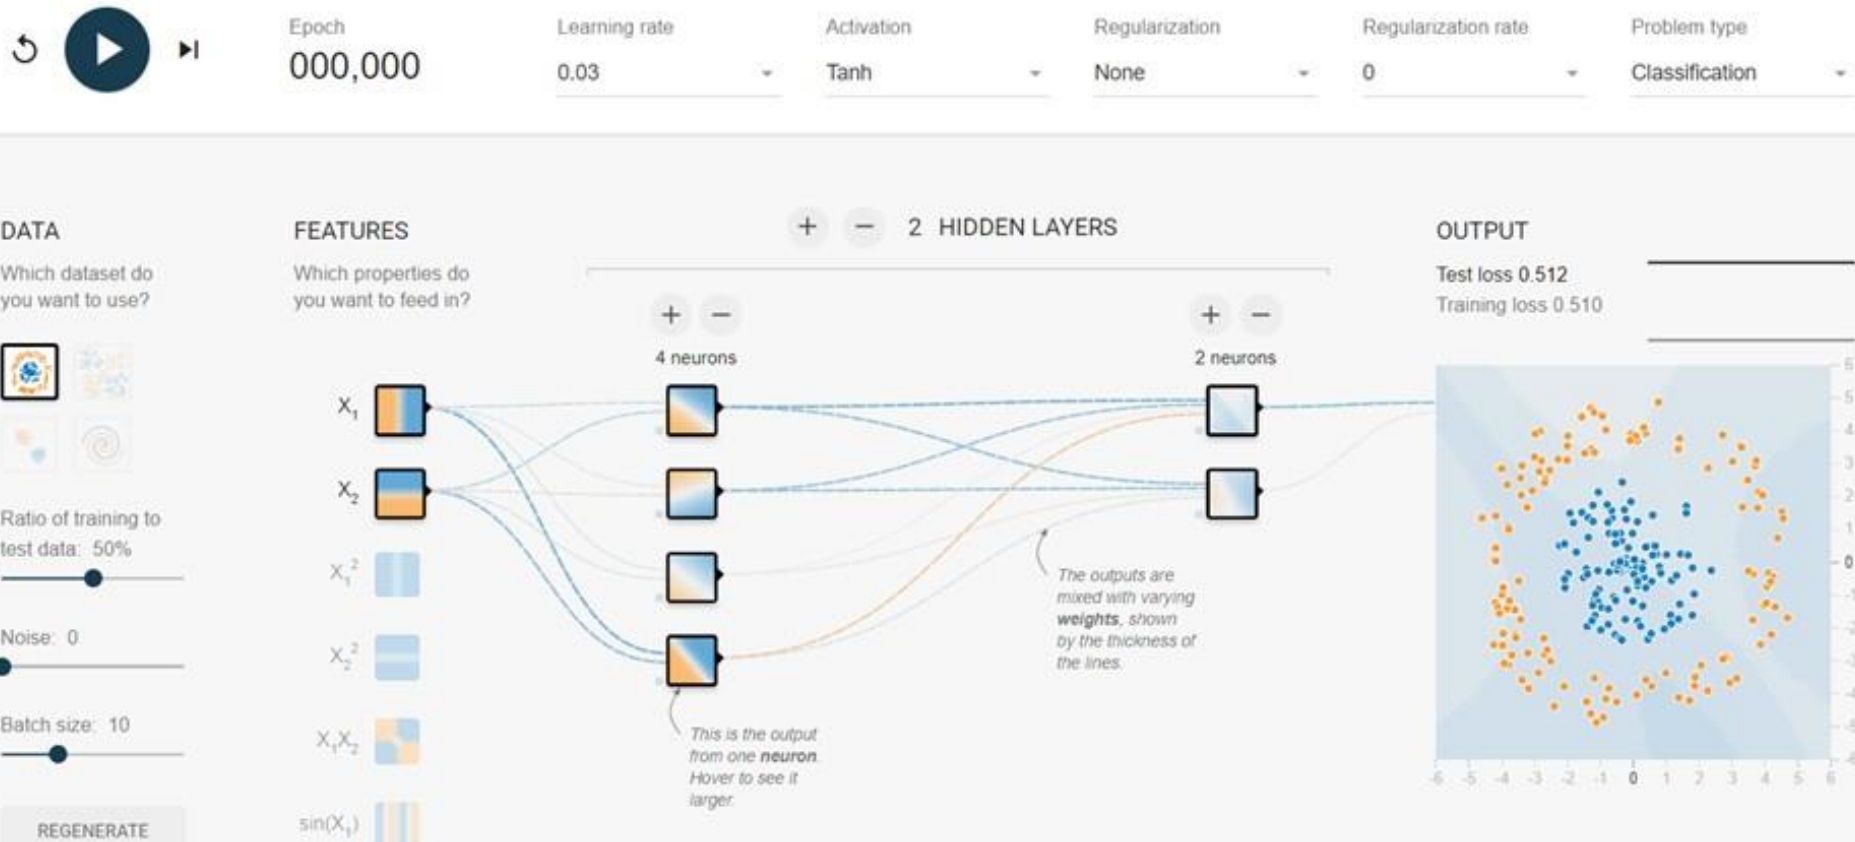
- Efficiency
- Memory Usage
- Coding Language Supporting
- Debugging Information
- Documentation
- Organization & Community

# Tensorflow

- From 
- All about computation graphs
- Easy visualizations (TensorBoard)
- Multi-GPU and multi-node training
- Easy deploying
- Have a try: <http://playground.tensorflow.org/>

# Tensorflow: Playground


Tinker With a **Neural Network** Right Here in Your Browser.  
Don't Worry, You Can't Break It. We Promise.



# Tensorflow: Version

- For this course, we are using Tensorflow 1.13
- Why we don't use Tensorflow 2.0?
  - Eager Execution
  - Unstable APIs

# Tensorflow: Computational Graphs



```
# Basic computational graph
import numpy as np
np.random.seed(0)
import tensorflow as tf

N, D = 3, 4


x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
z = tf.placeholder(tf.float32)

a = x * y
b = a + z
c = tf.reduce_sum(b)

grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])

with tf.Session() as sess:
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }
    out = sess.run([c, grad_x, grad_y, grad_z],
                    feed_dict=values)
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```

# Tensorflow: Computational Graphs



First, **define** the computational graph

```
# Basic computational graph
import numpy as np
np.random.seed(0)
import tensorflow as tf

N, D = 3, 4

x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
z = tf.placeholder(tf.float32)


a = x * y
b = a + z
c = tf.reduce_sum(b)

grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])

with tf.Session() as sess:
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }
    out = sess.run([c, grad_x, grad_y, grad_z],
                    feed_dict=values)
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```



# Tensorflow: Computational Graphs



First, **define** the computational graph

Second, **run** the graph many times

```
# Basic computational graph
```

```
import numpy as np
np.random.seed(0)
import tensorflow as tf
```

```
N, D = 3, 4
```

```
x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
z = tf.placeholder(tf.float32)
```

```
a = x * y
b = a + z
c = tf.reduce_sum(b)
```

```
grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])
```


```
with tf.Session() as sess:
```

```
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }
```

```
    out = sess.run([c, grad_x, grad_y, grad_z],
                    feed_dict=values)
```

```
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```

# Tensorflow: Computational Graphs



Define symbolic placeholders: inputs to the graph

```
# Basic computational graph
import numpy as np
np.random.seed(0)
import tensorflow as tf

N, D = 3, 4


x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
z = tf.placeholder(tf.float32)

a = x * y
b = a + z
c = tf.reduce_sum(b)

grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])

with tf.Session() as sess:
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }
    out = sess.run([c, grad_x, grad_y, grad_z],
                    feed_dict=values)
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```

# Tensorflow: Computational Graphs



Create forward  
computational graph →

```
# Basic computational graph
import numpy as np
np.random.seed(0)
import tensorflow as tf

N, D = 3, 4


x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
z = tf.placeholder(tf.float32)

a = x * y
b = a + z
c = tf.reduce_sum(b)

grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])

with tf.Session() as sess:
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }
    out = sess.run([c, grad_x, grad_y, grad_z],
                    feed_dict=values)
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```

# Tensorflow: Computational Graphs



Compute gradients →

```
# Basic computational graph
import numpy as np
np.random.seed(0)
import tensorflow as tf

N, D = 3, 4


x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
z = tf.placeholder(tf.float32)

a = x * y
b = a + z
c = tf.reduce_sum(b)

grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])

with tf.Session() as sess:
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }
    out = sess.run([c, grad_x, grad_y, grad_z],
                    feed_dict=values)
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```

# Tensorflow: Computational Graphs



Create **session** to run the graph

```
# Basic computational graph
import numpy as np
np.random.seed(0)
import tensorflow as tf

N, D = 3, 4

x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
z = tf.placeholder(tf.float32)


a = x * y
b = a + z
c = tf.reduce_sum(b)

grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])

with tf.Session() as sess:
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }
    out = sess.run([c, grad_x, grad_y, grad_z],
                    feed_dict=values)
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```



# Tensorflow: Computational Graphs



Create numpy arrays  
that will fill in the  
placeholders above

```
# Basic computational graph
import numpy as np
np.random.seed(0)
import tensorflow as tf

N, D = 3, 4


x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
z = tf.placeholder(tf.float32)

a = x * y
b = a + z
c = tf.reduce_sum(b)

grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])

with tf.Session() as sess:
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }
    out = sess.run([c, grad_x, grad_y, grad_z],
                    feed_dict=values)
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```

# Tensorflow: Computational Graphs



Run the graph: feed in the numpy arrays and get c\_val and gradients

```
# Basic computational graph
import numpy as np
np.random.seed(0)
import tensorflow as tf

N, D = 3, 4

x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
z = tf.placeholder(tf.float32)

a = x * y
b = a + z
c = tf.reduce_sum(b)

grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])

with tf.Session() as sess:
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }

    out = sess.run([c, grad_x, grad_y, grad_z],
                    feed_dict=values)
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```

# Tensorflow: About Tensor

- A+B Problem

```
import tensorflow as tf

a = tf.placeholder(tf.float32)
b = tf.placeholder(tf.float32)
c = tf.add(a, b)

with tf.Session() as session:
    result = session.run(c, {a:1, b:2})
    print(result) # 3
    result = session.run(c, {a:[5,2,1], b:[3,6,5]})
    print(result) # [8., 8., 6.]
```



# Tensorflow: About Tensor

- Everything is Tensor

```
import tensorflow as tf

a = tf.placeholder(tf.float32)
b = tf.placeholder(tf.float32)
c = tf.add(a, b)

print(a) # Tensor("Placeholder:0", dtype=float32)
print(b) # Tensor("Placeholder_1:0", dtype=float32)
print(c) # Tensor("add:0", dtype=float32)
```

# Tensorflow: About Tensor

- About Placeholder
  - 1.dtype, 2.shape, 3.name

```
a = tf.placeholder(tf.float32, shape=[3,5])
print(a)
# Tensor("Placeholder:0",shape=(3,5),dtype=float32)

b = tf.placeholder(tf.int32,
                  shape=[3,None,2], name='qq')
print(b)
# Tensor("qq:0",shape=(3,?,2),dtype=int32)
```


# Tensorflow: Work on MNIST

- A handwritten digits classification task.



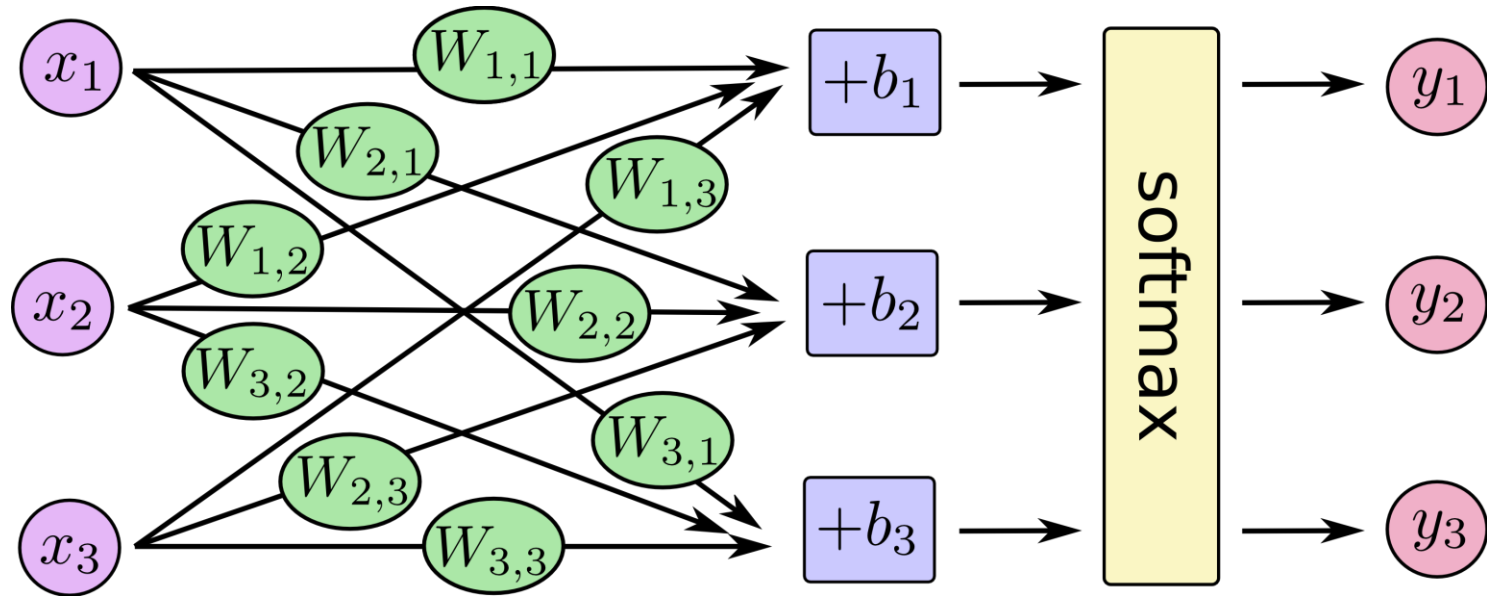
# Tensorflow: Work on MNIST

- A simple model
  - inputs:  $x$ , outputs:  $y$
  - parameters of model:  $W$ ,  $b$



# Tensorflow: Work on MNIST

- A simple model
  - inputs:  $x$ , outputs:  $y$
  - parameters of model:  $W$ ,  $b$



# Tensorflow: Work on MNIST

- A simple model
  - inputs:  $x$ , outputs:  $y$
  - parameters of model:  $W$ ,  $b$

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix} = \text{softmax} \left( \begin{bmatrix} W_{1,1}x_1 + W_{1,2}x_2 + W_{1,3}x_3 + b_1 \\ W_{2,1}x_1 + W_{2,2}x_2 + W_{2,3}x_3 + b_2 \\ W_{3,1}x_1 + W_{3,2}x_2 + W_{3,3}x_3 + b_3 \end{bmatrix} \right)$$

# Tensorflow: Work on MNIST

- A simple model
  - inputs:  $x$ , outputs:  $y$
  - parameters of model:  $W$ ,  $b$

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix} = \text{softmax} \left( \begin{bmatrix} W_{1,1} & W_{1,2} & W_{1,3} \\ W_{2,1} & W_{2,2} & W_{2,3} \\ W_{3,1} & W_{3,2} & W_{3,3} \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} + \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix} \right)$$

# Tensorflow: Work on MNIST

- A simple model
  - inputs:  $x$ , outputs:  $y$
  - parameters of model:  $W$ ,  $b$

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix} = \text{softmax} \left( \begin{bmatrix} W_{1,1} & W_{1,2} & W_{1,3} \\ W_{2,1} & W_{2,2} & W_{2,3} \\ W_{3,1} & W_{3,2} & W_{3,3} \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} + \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix} \right)$$

$$y = \text{softmax}(Wx + b)$$



# Tensorflow: Work on MNIST

- Build the graph

```
import tensorflow as tf

x = tf.placeholder(tf.float32, [None, 784], name='x')
# shape of x is (?,784)
W = tf.Variable(tf.zeros([784, 10]))
b = tf.Variable(tf.zeros([10]))
y = tf.matmul(x, W) + b
# shape of y is (?,10)
z = tf.argmax(y, 1, name='z')
```

# Tensorflow: Work on MNIST

- Compute the loss & gradients

```
y_hat = tf.placeholder(tf.float32, [None, 10])
loss = tf.nn.softmax_cross_entropy_with_logits(
    labels=y_hat, logits=y)
# shape of loss is (?)
loss_avg = tf.reduce_mean(loss)
# shape of loss_avg is ()

opt = tf.train.GradientDescentOptimizer(0.5)
gradients = opt.compute_gradients(loss_avg)
train_step = opt.apply_gradients(gradients)
```

# Tensorflow: Work on MNIST

- Training & Saving

```
mnist = input_data.read_data_sets(DIR, one_hot=True)
with tf.Session() as sess:
    tf.global_variables_initializer().run()
    for _ in range(1000):
        xs, ys = mnist.train.next_batch(100)
        sess.run(train_step, {x: xs, y_hat: ys})

saver = tf.train.Saver(tf.global_variables())
saver.save(sess, 'train/train')
```

# Tensorflow: Work on MNIST

- Restoring & Testing

```
mnist = input_data.read_data_sets(DIR, one_hot=True)
saver = tf.train.import_meta_graph('train/train.meta')
with tf.Session() as sess:
    saver.restore(sess, 'train/train')
    result = sess.run('z:0',{'x:0':mnist.test.images})
```

# Tensorflow: More Details

- Variable Scope

```
with tf.variable_scope('scope1'):
    v1 = tf.get_variable('v', shape=[23])
with tf.variable_scope('scope2'):
    v2 = tf.get_variable('v', shape=[23])

for item in tf.global_variables():
    print((item.name, item.get_shape()))

# (u'scope1/v:0', TensorShape([Dimension(23)]))
# (u'scope2/v:0', TensorShape([Dimension(23)]))
```

# Tensorflow: More Details

- Reuse
  - bad case

```
with tf.variable_scope('scope1'):  
    v1 = tf.get_variable('v', shape=[23])  
with tf.variable_scope('scope1'):  
    v2 = tf.get_variable('v', shape=[23])
```

ValueError: Variable scope1/v already exists

# Tensorflow: More Details

- Reuse
  - good case

```
with tf.variable_scope('scope1'):  
    v1 = tf.get_variable('v', shape=[23])  
with tf.variable_scope('scope1', reuse=True):  
    v2 = tf.get_variable('v', shape=[23])
```

v2 and v1 share the same parameter

# Tensorflow: More Details

- Trainable Variable

```
with tf.variable_scope('scope1'):
    v1 = tf.get_variable('v', shape=[23])

with tf.variable_scope('scope2'):
    v2 = tf.get_variable('v', shape=[23],
                        trainable=False)

for item in tf.trainable_variables():
    print((item.name, item.get_shape()))

# (u'scope1/v:0', TensorShape([Dimension(23)]))
```



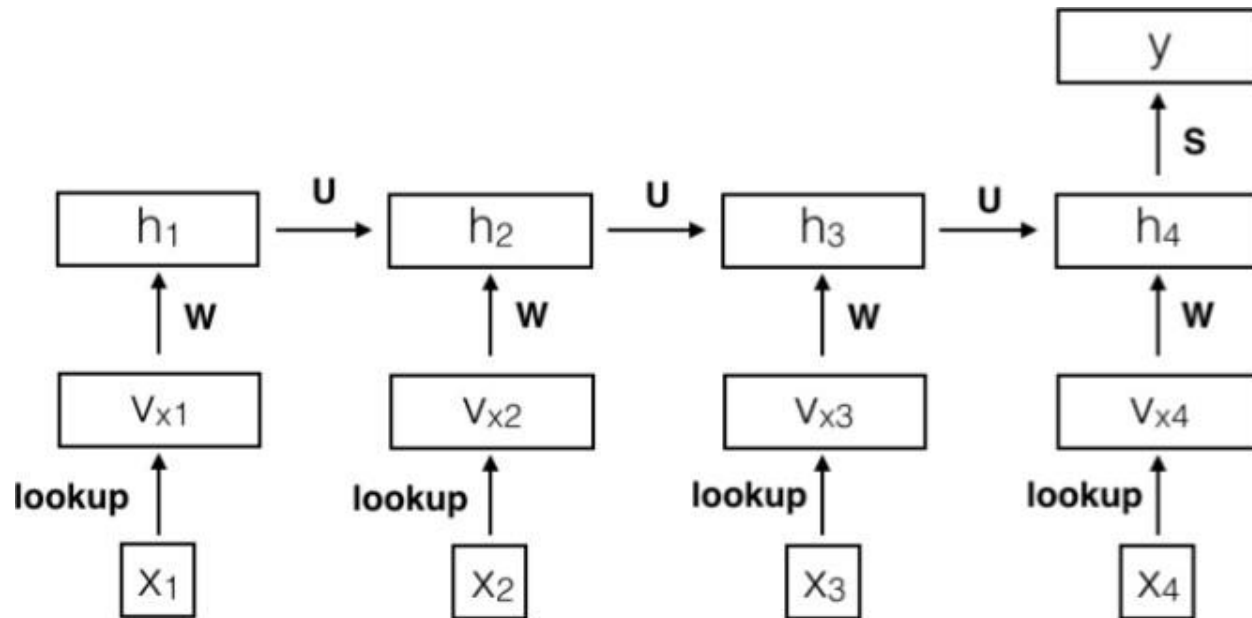
# Tensorflow: More Details

- Print all trainable variables for debugging

```
embed:0: (40000, 100)
encoder/multi_rnn_cell/cell_0/gru_cell/gates/weights:0: (1124, 2048)
encoder/multi_rnn_cell/cell_0/gru_cell/gates/biases:0: (2048,)
encoder/multi_rnn_cell/cell_0/gru_cell/candidate/weights:0: (1124, 1024)
encoder/multi_rnn_cell/cell_0/gru_cell/candidate/biases:0: (1024,)
encoder/multi_rnn_cell/cell_1/gru_cell/gates/weights:0: (2048, 2048)
encoder/multi_rnn_cell/cell_1/gru_cell/gates/biases:0: (2048,)
encoder/multi_rnn_cell/cell_1/gru_cell/candidate/weights:0: (2048, 1024)
encoder/multi_rnn_cell/cell_1/gru_cell/candidate/biases:0: (1024,)
encoder/multi_rnn_cell/cell_2/gru_cell/gates/weights:0: (2048, 2048)
encoder/multi_rnn_cell/cell_2/gru_cell/gates/biases:0: (2048,)
encoder/multi_rnn_cell/cell_2/gru_cell/candidate/weights:0: (2048, 1024)
encoder/multi_rnn_cell/cell_2/gru_cell/candidate/biases:0: (1024,)
encoder/multi_rnn_cell/cell_3/gru_cell/gates/weights:0: (2048, 2048)
encoder/multi_rnn_cell/cell_3/gru_cell/gates/biases:0: (2048,)
encoder/multi_rnn_cell/cell_3/gru_cell/candidate/weights:0: (2048, 1024)
encoder/multi_rnn_cell/cell_3/gru_cell/candidate/biases:0: (1024,)
attention_keys/weights:0: (1024, 1024)
decoder/multi_rnn_cell/cell_0/gru_cell/gates/weights:0: (2148, 2048)
decoder/multi_rnn_cell/cell_0/gru_cell/gates/biases:0: (2048,)
decoder/multi_rnn_cell/cell_0/gru_cell/candidate/weights:0: (2148, 1024)
decoder/multi_rnn_cell/cell_0/gru_cell/candidate/biases:0: (1024,)
decoder/multi_rnn_cell/cell_1/gru_cell/gates/weights:0: (2048, 2048)
decoder/multi_rnn_cell/cell_1/gru_cell/gates/biases:0: (2048,)
decoder/multi_rnn_cell/cell_1/gru_cell/candidate/weights:0: (2048, 1024)
decoder/multi_rnn_cell/cell_1/gru_cell/candidate/biases:0: (1024,)
decoder/multi_rnn_cell/cell_2/gru_cell/gates/weights:0: (2048, 2048)
decoder/multi_rnn_cell/cell_2/gru_cell/gates/biases:0: (2048,)
decoder/multi_rnn_cell/cell_2/gru_cell/candidate/weights:0: (2048, 1024)
decoder/multi_rnn_cell/cell_2/gru_cell/candidate/biases:0: (1024,)
decoder/multi_rnn_cell/cell_3/gru_cell/gates/weights:0: (2048, 2048)
decoder/multi_rnn_cell/cell_3/gru_cell/gates/biases:0: (2048,)
decoder/multi_rnn_cell/cell_3/gru_cell/candidate/weights:0: (2048, 1024)
decoder/multi_rnn_cell/cell_3/gru_cell/candidate/biases:0: (1024,)
attention_construct/weights:0: (2048, 1024)
decoder/output_projection/weights:0: (1024, 40000)
decoder/output_projection/biases:0: (40000,)
```

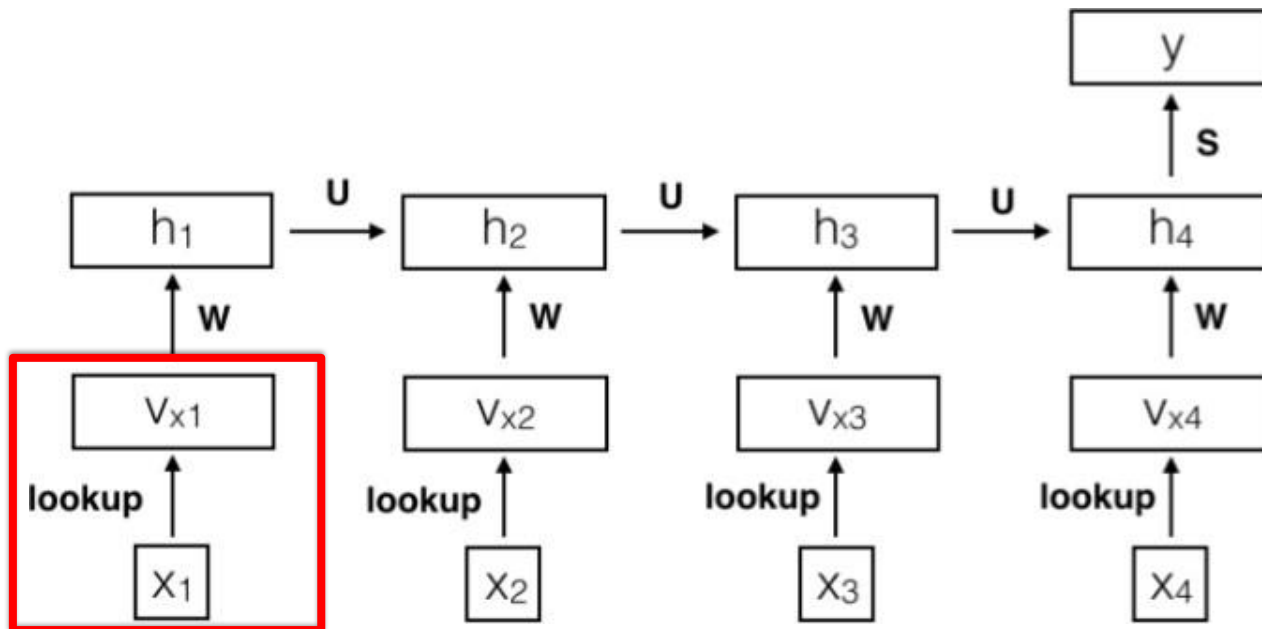
# Tensorflow: About RNN

- RNN for Sentence Classification



# Tensorflow: About RNN

- RNN for Sentence Classification
  - Word Embedding



# Tensorflow: About RNN

- Word Embedding

```
# initialize randomly  
embed = tf.get_variable('embed', [vocab_size, d],  
                        tf.float32)
```

```
# initialize by pre-trained word vectors  
embed = tf.get_variable('embed', dtype=tf.float32,  
                        initializer=init_embed)
```

```
# using pre-trained word vectors without tuning  
embed = tf.get_variable('embed', dtype=tf.float32,  
                        initializer=init_embed, trainable=False)
```

# Tensorflow: About RNN

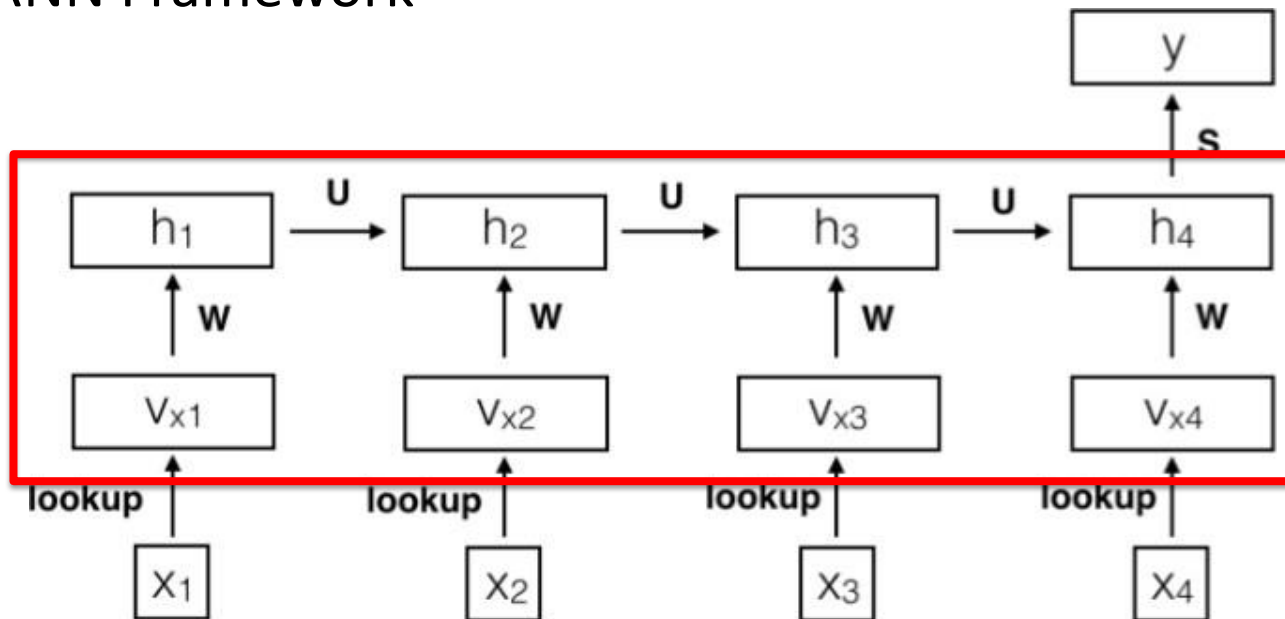
- Embedding Lookup

```
# initialize randomly  
embed = tf.get_variable('embed', [vocab_size, d],  
                        tf.float32)
```

```
x = tf.placeholder(tf.int32, shape=[None, max_step])  
v = tf.nn.embedding_lookup(embed, x)  
# the shape of v is (?,max_step,d)
```

# Tensorflow: About RNN

- RNN for Sentence Classification
  - RNN Framework



# Tensorflow: About RNN

- A Simple RNN

```
h = tf.zeros((1, d))
for step in range(max_step):
    with tf.variable_scope('rnn', reuse=step>0):
        with tf.variable_scope('W'):
            parts_w = layers.linear(v[:,step,:], d)
        with tf.variable_scope('U'):
            parts_u = layers.linear(h, d)
        h = tf.nn.tanh(parts_w+parts_u)

with tf.variable_scope('classifier'):
    y = layers.linear(h, 5)

# 'embed:0': (10000, 300)
# 'rnn/W/fully_connected/weights:0': (300, 300)
# 'rnn/W/fully_connected/biases:0': (300)
# 'rnn/U/fully_connected/weights:0': (300, 300)
# 'rnn/U/fully_connected/biases:0': (300)
# 'classifier/fully_connected/weights:0': (300, 5)
# 'classifier/fully_connected/biases:0': (5)
```

# Tensorflow: About RNN

- Using RNNCell Module

```
from tensorflow.contrib.rnn import BasicRNNCell

cell = BasicRNNCell(d)
h = cell.zero_state(batch_size, tf.float32)

for step in range(max_step):
    with tf.variable_scope('rnn', reuse=step>0):
        _, h = cell(v[:,step,:], h)

# 'rnn/basic_rnn_cell/weights:0': (600, 300)
# 'rnn/basic_rnn_cell/biases:0': (300)
```



# Tensorflow: About RNN

- Using LSTMCell or GRUCell Module

```
from tensorflow.contrib.rnn import GRUCell, LSTMCell

cell = LSTMCell(d)
# 'rnn/lstm_cell/weights:0': (600, 1200)
# 'rnn/lstm_cell/biases:0': (1200)

cell = GRUCell(d)
# 'rnn/gru_cell/gates/weights:0': (600, 600)
# 'rnn/gru_cell/gates/biases:0': (600)
# 'rnn/gru_cell/candidate/weights:0': (600, 300)
# 'rnn/gru_cell/candidate/biases:0': (300)
```

# Tensorflow: About RNN

- Using DynamicRNN

```
x = tf.placeholder(tf.int32, shape=[None, None])
x_len = tf.placeholder(tf.int32, shape=[None])
# an example for x and x_len
# x = [[2, 4, 6, 0], [3, 5, 6, 1], [2, 4, 0, 0]]
# x_len = [3, 4, 2]

v = tf.nn.embedding_lookup(embed, x)
# the shape of v is (?, ?, d)
```

```
from tensorflow.nn import dynamic_rnn

cell = BasicRNNCell(d)
_, h = dynamic_rnn(cell, v, x_len, dtype=tf.float32)

# 'rnn/basic_rnn_cell/weights:0': (600, 300)
# 'rnn/basic_rnn_cell/biases:0': (300)
```

# Tensorflow: How to Debug

- tf.Print

```
import tensorflow as tf

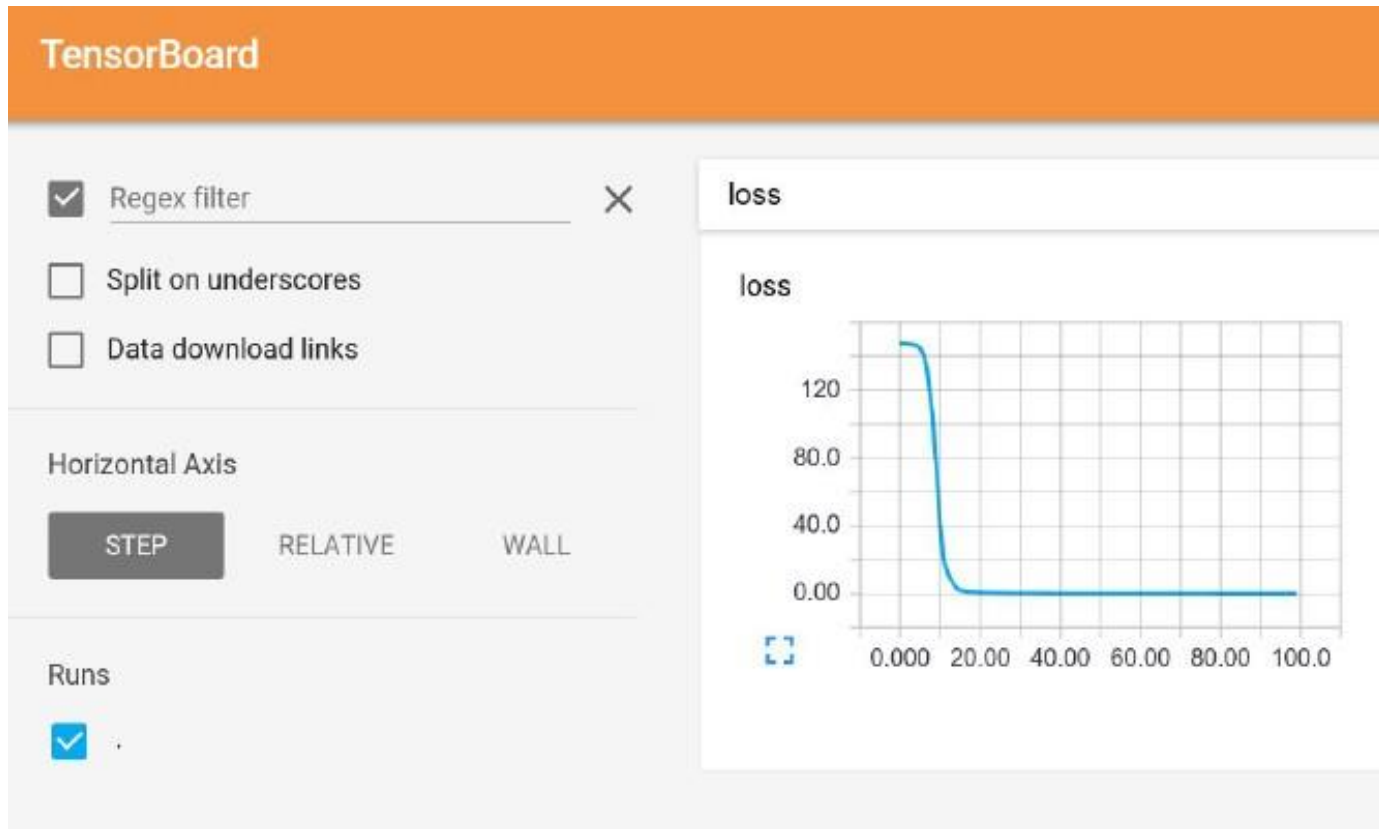
a = tf.placeholder(tf.float32)
b = tf.placeholder(tf.float32)
c = tf.add(a, b)
c = tf.Print(c, [tf.shape(c)], summarize=10)

with tf.Session() as session:
    result = session.run(c, {a:[1, 2, 3], b:[4, 5, 6]})
    #I tensorflow/core/kernels/logging_ops.cc:79] [3]


    result = session.run(c, {a:[[1,2],[3,4]], b:[[4,5], [4,6]]})
    #I tensorflow/core/kernels/logging_ops.cc:79] [2 2]
```

# Tensorflow: Tensorboard

- Add logging to code to record loss, stats, etc.



# Tensorflow: Distributed Version



Split one graph  
over multiple  
machines!

<https://tensorflow.google.cn/deploy/distributed>

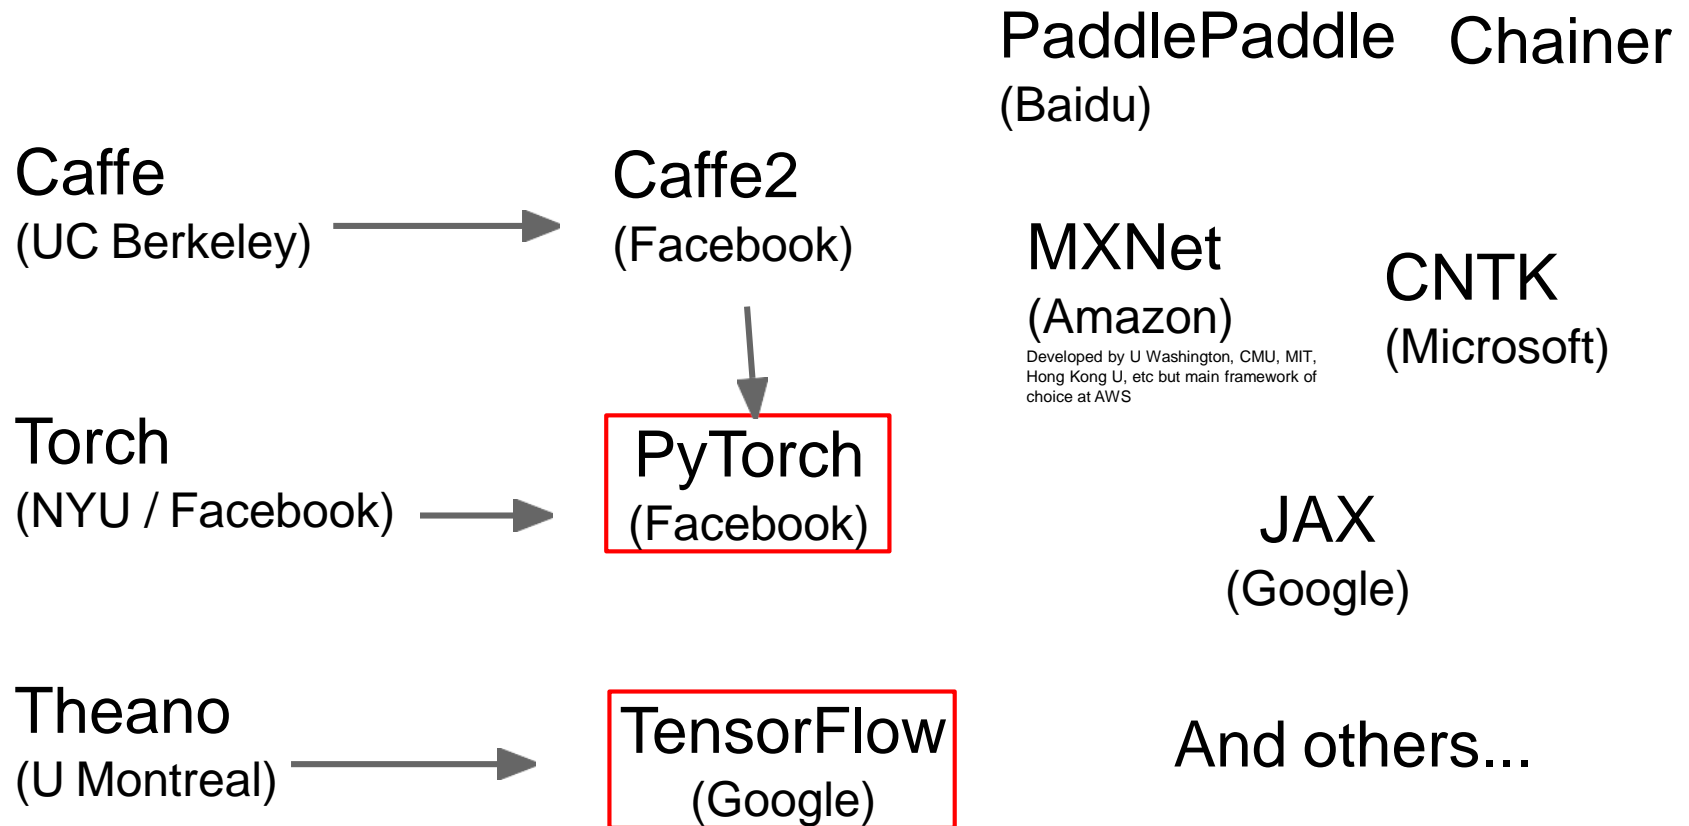
# Tensorflow: Summary

- Placeholder(dtype,shape,name)
- Variable(dtype,shape,name,scope,reuse,trainable)
- Tensor(dtype,shape)
- Operation, Graph, Session
- Gradient, Optimizer
- Save, Restore, Debug (Tensorboard)
- Distributed Version

# Tensorflow: High-level Wrappers

- `tf.layers`  
([https://tensorflow.google.cn/api\\_docs/python/tf/layers](https://tensorflow.google.cn/api_docs/python/tf/layers))
- `tf.contrib.layers`  
([https://tensorflow.google.cn/api\\_docs/python/tf/contrib/layers](https://tensorflow.google.cn/api_docs/python/tf/contrib/layers))
- `tf.estimator`  
([https://tensorflow.google.cn/api\\_docs/python/tf/estimator](https://tensorflow.google.cn/api_docs/python/tf/estimator))
- `tf.keras`  
([https://tensorflow.google.cn/api\\_docs/python/tf/keras](https://tensorflow.google.cn/api_docs/python/tf/keras))
- Sonnet (<https://github.com/deepmind/sonnet>)
- TFLearn (<http://tflearn.org/>)
- TensorLayer (<https://tensorlayer.readthedocs.io/en/latest/>)
- ZhuSuan (<https://zhusuan.readthedocs.io/en/latest/>)


# A zoo of frameworks!



We'll focus on these



# PyTorch: An example




```
import torch

N, D = 3, 4
x = torch.randn(N, D)
y = torch.randn(N, D)
z = torch.randn(N, D)

a = x * y
b = a + z
c = torch.sum(b)
```

Looks exactly like numpy!

# PyTorch: An example



```
import torch
```

```
N, D = 3, 4
```

```
x = torch.randn(N, D, requires_grad=True)
```

```
y = torch.randn(N, D)
```

```
z = torch.randn(N, D)
```

```
a = x * y
```

```
b = a + z
```


```
c = torch.sum(b)
```

```
c.backward()
```

```
print(x.grad)
```

PyTorch handles gradients for us!

# PyTorch: An example



```
import torch
device = 'cuda:0'
N, D = 3, 4
x = torch.randn(N, D, requires_grad=True,
                 device=device)
y = torch.randn(N, D, device=device)
z = torch.randn(N, D, device=device)

a = x * y
b = a + z
c = torch.sum(b)

c.backward()
print(x.grad)
```

Trivial to run on GPU  
– just construct arrays on a different device!

# PyTorch: Version

- For this slides, we are using PyTorch 1.0
- Be careful if you are looking at older PyTorch code!

# PyTorch: Fundamental Concepts

- **Tensor:** Like a numpy array, but can run on GPU
- **Autograd:** Package for building computational graphs out of Tensors, and automatically computing gradients
- **Module:** A neural network layer; may store state or learnable weights

# PyTorch: Autograd

We will not want gradients  
(of loss) with respect to  
data

Do want gradients with  
respect to weights

Operations on Tensors with  
`requires_grad=True` cause  
PyTorch to build a  
computational graph

```
import torch

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)
w1 = torch.randn(D_in, H, requires_grad=True)
w2 = torch.randn(H, D_out, requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    loss.backward()

    with torch.no_grad():
        w1 -= learning_rate * w1.grad
        w2 -= learning_rate * w2.grad
        w1.grad.zero_()
        w2.grad.zero_()
```

# PyTorch: Autograd

Forward pass looks exactly the same as before, but we don't need to track intermediate values - PyTorch keeps track of them for us in the graph

```
import torch

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)
w1 = torch.randn(D_in, H, requires_grad=True)
w2 = torch.randn(H, D_out, requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    loss.backward()

    with torch.no_grad():
        w1 -= learning_rate * w1.grad
        w2 -= learning_rate * w2.grad
        w1.grad.zero_()
        w2.grad.zero_()
```

# PyTorch: Autograd

Compute gradient of loss  
with respect to w1 and  
w2

```
import torch

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)
w1 = torch.randn(D_in, H, requires_grad=True)
w2 = torch.randn(H, D_out, requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    loss.backward()

    with torch.no_grad():
        w1 -= learning_rate * w1.grad
        w2 -= learning_rate * w2.grad
        w1.grad.zero_()
        w2.grad.zero_()
```



# PyTorch: Autograd

```
import torch

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)
w1 = torch.randn(D_in, H, requires_grad=True)
w2 = torch.randn(H, D_out, requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    loss.backward()

    with torch.no_grad():
        w1 -= learning_rate * w1.grad
        w2 -= learning_rate * w2.grad
        w1.grad.zero_()
        w2.grad.zero_()
```

Make gradient step on weights, then zero them. `Torch.no_grad` means “don’t build a computational graph for this part”

# PyTorch: Autograd

Clear the gradients, waiting for next loop

PyTorch methods that end in underscore modify the Tensor in-place; methods that don't return a new Tensor

```
import torch

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)
w1 = torch.randn(D_in, H, requires_grad=True)
w2 = torch.randn(H, D_out, requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    loss.backward()

    with torch.no_grad():
        w1 -= learning_rate * w1.grad
        w2 -= learning_rate * w2.grad
        w1.grad.zero_()
        w2.grad.zero_()
```

# PyTorch: nn

Higher-level wrapper for working with neural nets

Use this! It will make your life easier

```
import torch

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)

model = torch.nn.Sequential(
    torch.nn.Linear(D_in, H),
    torch.nn.ReLU(),
    torch.nn.Linear(H, D_out))

learning_rate = 1e-2
for t in range(500):
    y_pred = model(x)
    loss = torch.nn.functional.mse_loss(y_pred, y)

    loss.backward()

    with torch.no_grad():
        for param in model.parameters():
            param -= learning_rate * param.grad
    model.zero_grad()
```

# PyTorch: optim

Use an **optimizer** for different update rules

After computing gradient:  
use optimizer to update  
params and zero gradient

```
import torch

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)

model = torch.nn.Sequential(
    torch.nn.Linear(D_in, H),
    torch.nn.ReLU(),
    torch.nn.Linear(H, D_out))

learning_rate = 1e-4
optimizer = torch.optim.Adam(model.parameters(),
                               lr=learning_rate)

for t in range(500):
    y_pred = model(x)
    loss = torch.nn.functional.mse_loss(y_pred, y)

    loss.backward()

    optimizer.step()
    optimizer.zero_grad()
```



# PyTorch: nn -- Define new Modules

A PyTorch **Module** is a neural net layer; it inputs and outputs Tensors

Modules can contain weights or other modules

You can define your own Modules using autograd!

```
import torch

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear1 = torch.nn.Linear(D_in, H)
        self.linear2 = torch.nn.Linear(H, D_out)

    def forward(self, x):
        h_relu = self.linear1(x).clamp(min=0)
        y_pred = self.linear2(h_relu)
        return y_pred

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)

model = TwoLayerNet(D_in, H, D_out)

optimizer = torch.optim.SGD(model.parameters(), lr=1e-4)
for t in range(500):
    y_pred = model(x)
    loss = torch.nn.functional.mse_loss(y_pred, y)

    loss.backward()
    optimizer.step()
    optimizer.zero_grad()
```

# PyTorch: Dynamic Computation Graphs

```
import torch

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)
w1 = torch.randn(D_in, H, requires_grad=True)
w2 = torch.randn(H, D_out, requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    loss.backward()
```

# PyTorch: Dynamic Computation Graphs

x

w1

w2

y

```
import torch
```

```
N, D_in, H, D_out = 64, 1000, 100, 10
```

```
x = torch.randn(N, D_in)
```

```
y = torch.randn(N, D_out)
```

```
w1 = torch.randn(D_in, H, requires_grad=True)
```

```
w2 = torch.randn(H, D_out, requires_grad=True)
```

```
learning_rate = 1e-6
```

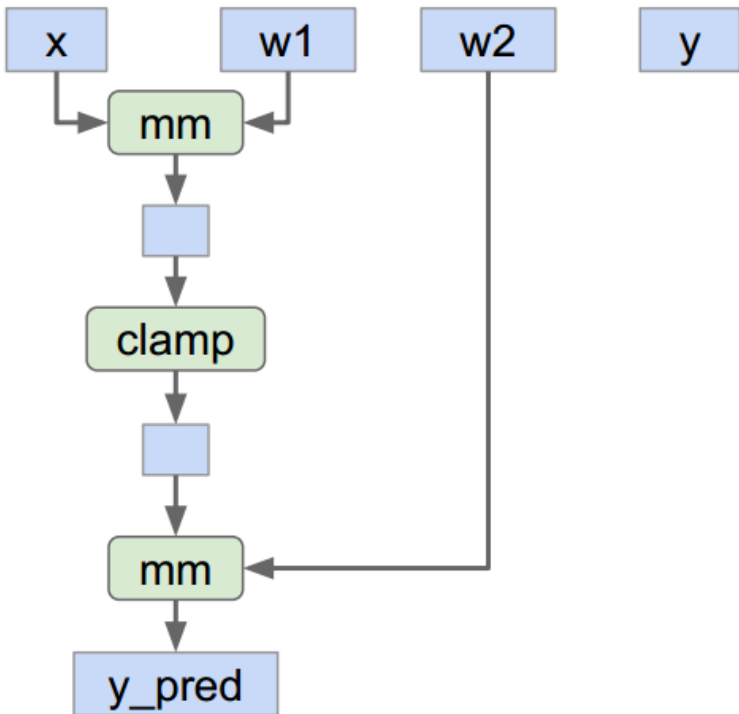
```
for t in range(500):
```

```
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
```

```
    loss = (y_pred - y).pow(2).sum()
```

```
    loss.backward()
```

# PyTorch: Dynamic Computation Graphs



```
import torch

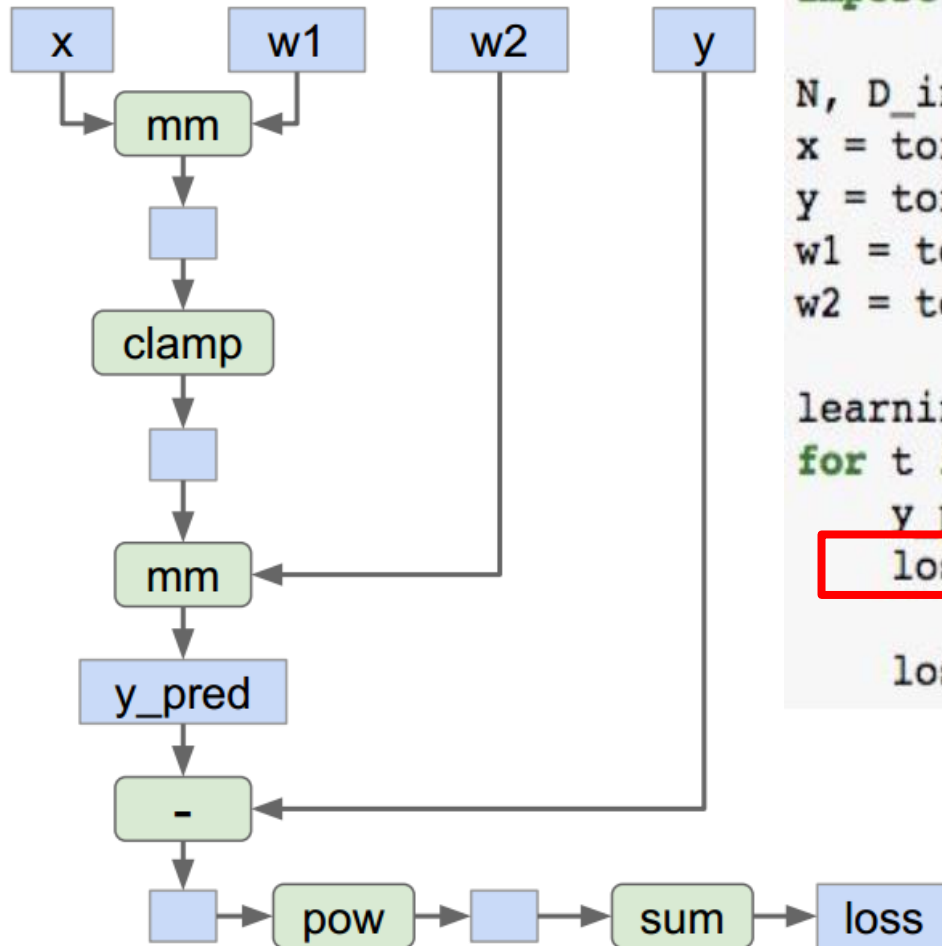
N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)
w1 = torch.randn(D_in, H, requires_grad=True)
w2 = torch.randn(H, D_out, requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    loss.backward()
```



# PyTorch: Dynamic Computation Graphs



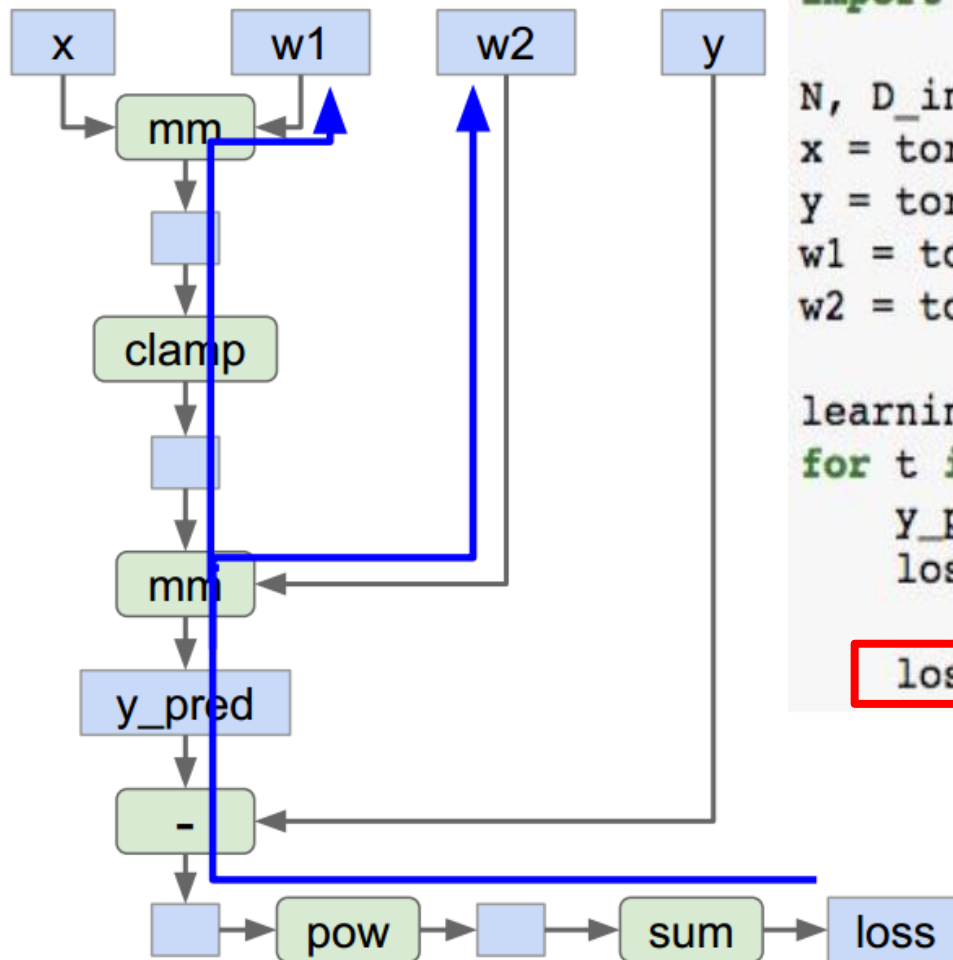
```
import torch

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)
w1 = torch.randn(D_in, H, requires_grad=True)
w2 = torch.randn(H, D_out, requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    loss.backward()
```

# PyTorch: Dynamic Computation Graphs



```
import torch

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)
w1 = torch.randn(D_in, H, requires_grad=True)
w2 = torch.randn(H, D_out, requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    loss.backward()
```

# PyTorch: Dynamic Computation Graphs

x

w1

w2

y

Throw away the graph,  
backprop path, and  
rebuild it from scratch on  
every iteration

```
import torch

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)
w1 = torch.randn(D_in, H, requires_grad=True)
w2 = torch.randn(H, D_out, requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    loss.backward()
```

# PyTorch: Dynamic Computation Graphs

**Building** the graph and  
**computing** the graph  
happen at the same time.

Seems inefficient,  
especially if we are  
building the same graph  
over and over again...

But useful if you are tackle  
with varied-length data

```
import torch

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in)
y = torch.randn(N, D_out)
w1 = torch.randn(D_in, H, requires_grad=True)
w2 = torch.randn(H, D_out, requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    loss.backward()
```

# Pytorch: TensorboardX

- A python wrapper around Tensorflow's web-based visualization tool.
  - `pip install tensorboardx`

# Static vs Dynamic Graphs

**TensorFlow:** Build graph once, then run many times (**static**)

```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.Variable(tf.random_normal((D, H)))
w2 = tf.Variable(tf.random_normal((H, D)))

h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))
grad_w1, grad_w2 = tf.gradients(loss, [w1, w2])

learning_rate = 1e-5
new_w1 = w1.assign(w1 - learning_rate * grad_w1)
new_w2 = w2.assign(w2 - learning_rate * grad_w2)
updates = tf.group(new_w1, new_w2)

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    values = {x: np.random.randn(N, D),
              y: np.random.randn(N, D),}
    losses = []
    for t in range(50):
        loss_val, _ = sess.run([loss, updates],
                               feed_dict=values)
```

Build  
graph

Run each  
iteration

**PyTorch:** Each forward pass defines a new graph (**dynamic**)

```
import torch
from torch.autograd import Variable

N, D_in, H, D_out = 64, 1000, 100, 10
x = Variable(torch.randn(N, D_in), requires_grad=False)
y = Variable(torch.randn(N, D_out), requires_grad=False)
w1 = Variable(torch.randn(D_in, H), requires_grad=True)
w2 = Variable(torch.randn(H, D_out), requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    if w1.grad: w1.grad.data.zero_()
    if w2.grad: w2.grad.data.zero_()
    loss.backward()

    w1.data -= learning_rate * w1.grad.data
    w2.data -= learning_rate * w2.grad.data
```

New graph each iteration



# Dynamic TensorFlow: Eager Execution

TensorFlow 2.0 supports **eager execution** which allows dynamic graphs!

Convert input numpy arrays to TF **tensors**.  
Create weights as `tf.Variable`

Use `tf.GradientTape()` context to build **dynamic** computation graph.

All forward-pass operations in the contexts (including function calls) gets traced for computing gradient later.

```
N, D, H = 64, 1000, 100

x = tf.convert_to_tensor(np.random.randn(N, D), np.float32)
y = tf.convert_to_tensor(np.random.randn(N, D), np.float32)
w1 = tf.Variable(tf.random.uniform((D, H))) # weights
w2 = tf.Variable(tf.random.uniform((H, D))) # weights

with tf.GradientTape() as tape:
    h = tf.maximum(tf.matmul(x, w1), 0)
    y_pred = tf.matmul(h, w2)
    diff = y_pred - y
    loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))
gradients = tape.gradient(loss, [w1, w2])
```

`tape.gradient()` uses the traced computation graph to compute gradient for the weights

# Personal Advice

- **TensorFlow pre2.0:** safe bet for most projects, not perfect but huge community, wide usage. Better pair with high-level wrapper (Keras, Sonnet, etc)
- **Tensorflow 2.0:** still new
- **PyTorch:** Personal choice. Best for research. Lighter than tensorflow, easy to debug.



# Links

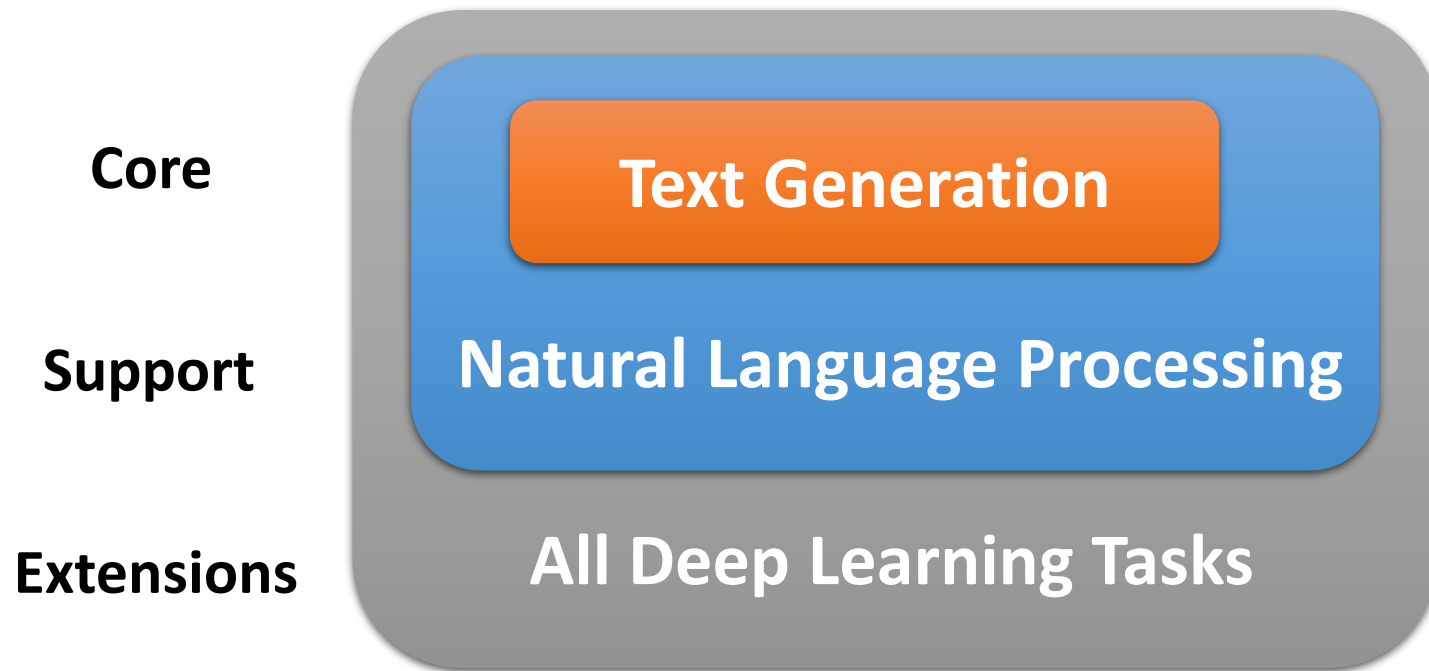
- <https://tensorflow.google.cn>
- <https://pytorch.org>
- <https://keras.io>

# Conversational Toolkits (CoTK)

HUANG Fei

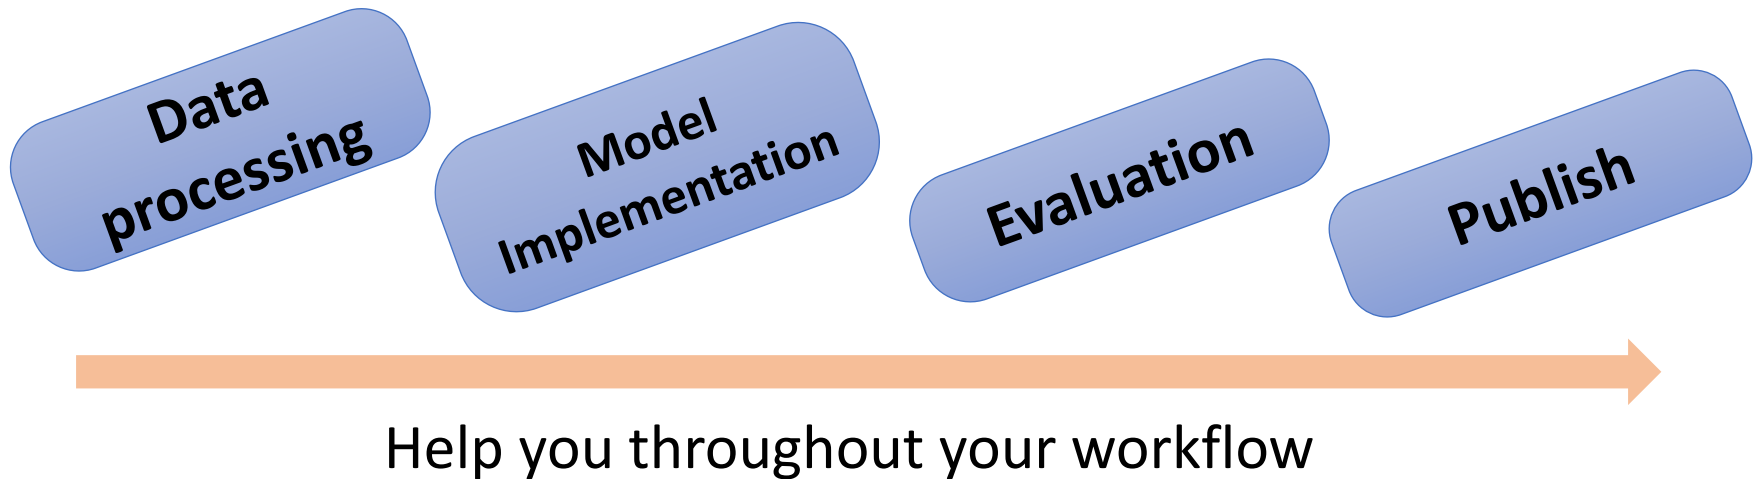
# What is CoTK

- A toolkit for deep learning researchers



Design Concept

# What can CoTK do



<b>Data processing</b>	Datasets and Preprocess	torchtext
<b>Model Implementation</b>	Baselines	Github
<b>Evaluation</b>	Metrics	NLTK
<b>Publish</b>	Reproduce Experiments	SACRED

# Advantages

- **A few lines** to start
- Light-weight and **framework independent**
- Predefined standard datasets
- Predefined baseline models
- Compare models **fairly**
- **Reproduce** your and others' experiments

# Overview

- An example: Implement a GRU LM
- Quick Start
  - Installation
  - Dataloader
  - Metrics
  - Publish Experiments
  - Reproduce Experiments
  - Predefined models
- Extending CoTK

# Quick Start

# Installation

- cmd: `pip install cotk`
- Github: `github.com/thu-coai/cotk`
- Homepage:  
<http://coai.cs.tsinghua.edu.cn/dialtk/cotk/>
- Tutorials:  
`https://thu-coai.github.io/cotk_docs/`



# Dataloader

- Automatically download online resources
  - `cotk.dataloader.MSCOCO("resources://MSCOCO_small")`
- Download from a url
  - `cotk.dataloader.MSCOCO("http://cotk-data.s3-ap-northeast-1.amazonaws.com/mscoco_small.zip#MSCOCO")`
- Import from local file
  - `cotk.dataloader.MSCOCO("./MSCOCO.zip#MSCOCO")`

**NAME @ SOURCE # Preprocessor**

# Dataloader

- Inspect vocabulary list
  - `dataloader.vocab_size`  
Vocabulary size: 2588
  - `dataloader.vocab_list[:10]`  
['<pad>', '<unk>', '<go>', '<eos>', '.', 'a', 'A', 'on', 'of', 'in']
- Convert between ids and strings
  - `dataloader.convert_tokens_to_ids(["<go>", "hello", "world", "<eos>"])`
  - `dataloader.convert_ids_to_tokens([2, 1379, 1897, 3])`

# Dataloader

- Iterative over batch
  - for data in dataloader.get\_batch("train", batch\_size=1):  
print(data)

**{'sent':**

```
array([[ 2, 181, 13, 26, 145, 177, 8, 22, 12, 5, 1, 1099, 4, 3]]),  
# <go> This is an old photo of people and a <unk> wagon.
```

**'sent\_allvocabs':**


```
array([[ 2, 181, 13, 26, 145, 177, 8, 22, 12, 5, 3755, 1099, 4, 3]]),  
# <go> This is an old photo of people and a horse-drawn wagon.
```

**'sent\_length':** array([14])

**}**

# Valid / Invalid / Unknown Vocabs


- Valid vocabs
  - From training set
  - Appear  $> \text{min\_vocab\_times}$
  - Model should read and generate
- Invalid vocabs
  - From train & dev & test set
  - Appear  $> \text{invalid\_vocab\_times}$ , but not valid
  - Model can optionally read or generate in test stage (like copyNet)



# Valid / Invalid / Unknown Vocabs

- Unknown vocabs

- Not valid vocabs or invalid vocabs
- Model can't know or generate



- Why ?

- Transfer models between different datasets
- Compare between models with the same allvocabs
- Compare between common generate model & copyNet

# Valid / Invalid / Unknown Vocabs

- How ?
  - Most models only care about valid vocabs
    - Transferable if using the same valid vocabs
  - Metrics care about valid & invalid vocabs
    - Perplexity -> smoothing <unk> to invalid vocabs
    - Bleu -> <unk> never matched
    - Comparable if all\_vocabs is the same

# Dataloader

- LanguageGeneration
  - MSCOCO
- SingleTurnDialog
  - OpenSubtitles
- BERTSingleTurnDialog
  - BERTOpenSubtitles
- MultiTurnDialog
  - UbuntuCorpus
- SentenceClassification
  - SST

# Metrics

- Metric Pipeline

```
metric = cotk.metric.xxMetric(dataloader)
metric.forward(...)
metric.close()
```

- MetricChain: Merge multiple metrics

```
metric = cotk.metric.MetricChain()
metric.add_metric(metricA)
metric.add_metric(metricB)
metric.forward(...)
metric.close()
```



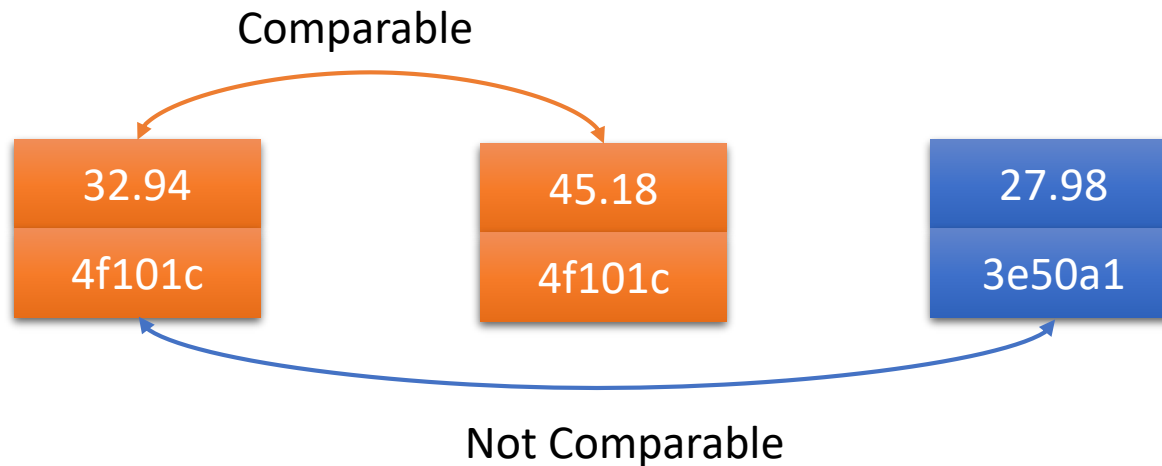
# Metrics

- Predefined metrics for given tasks

```
metric = dataloader.get_inference_metric(gen_key="gen")
metric.forward({
    "gen":
        [[2, 181, 13, 26, 145, 177, 8, 22, 12, 5, 3755, 1099, 4, 3],
         [2, 46, 145, 500, 1764, 207, 11, 5, 93, 7, 31, 4, 3]]
})
print(metric.close())
{
    'self-bleu': 0.0220,
    'self-bleu hashvalue': 'c206..',
    'fw-bleu': 0.383, 'bw-bleu': 0.0259, 'fw-bw-bleu': 0.0486
    'fw-bw-bleu hashvalue': '530d...',
}
```

# Metrics

- Hashvalue: Make sure we have the same
  - References
  - Settings
  - Version

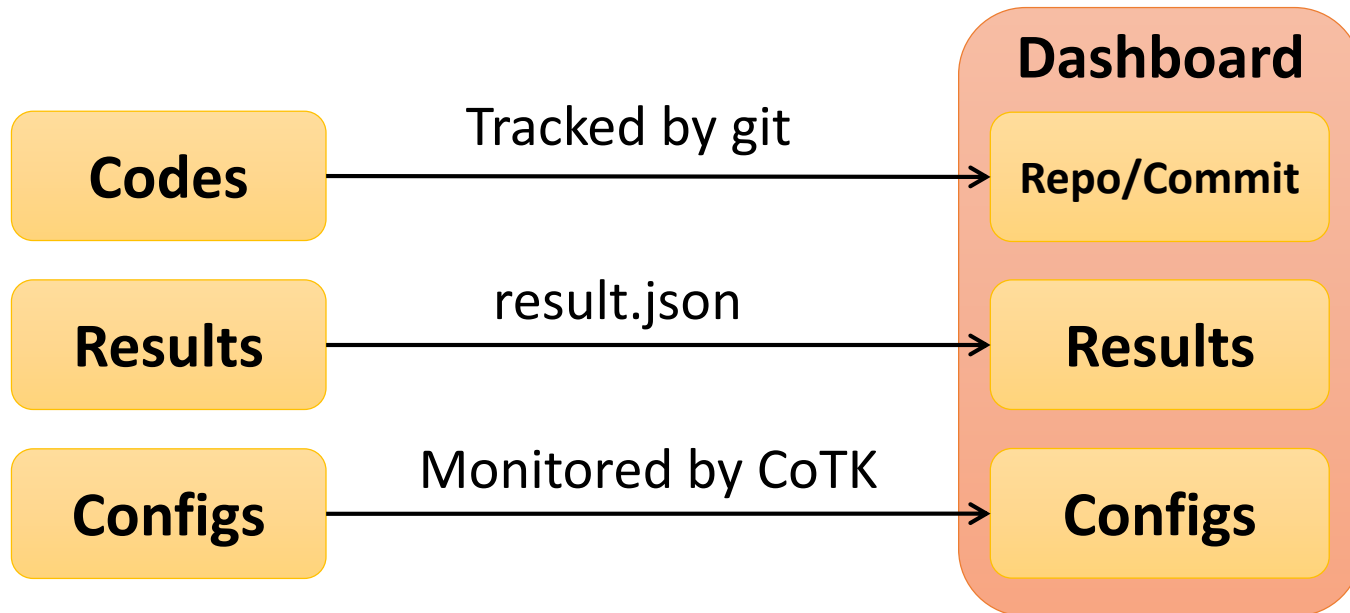


# Metrics

- PerplexityMetric
- BleuCorpusMetric
- SelfBleuCorpusMetric
- FwBwBleuCorpusMetric
- BleuPrecisionRecallMetric
- EmbSimilarityPrecisionRecallMetric
- NgramFwBwPerplexityMetric

# Publish Experiments

- Make a git repo
- Write your model
  - Output your result to result.json
- Use “cotk run” command



# Publish Experiments

Extra columns (Corresponding to keys in 'result', separated by commas):

perplexity,perplexity hashvalue,bleu,bleu hashvalue

Submit

Uploaded by

Github commit

file\_id

Search

Show 10 entries

ID <span>↑↓</span>	User <span>↑↓</span>	Github commit	Dataloader	perplexity <span>↑↓</span>	perplexity hashvalue
#5	hzhwcmhf	thu-coai/seq2seq-pytorch@aa1869	OpenSubtitles (resources://OpenSubtitles_small)	94.698	460f26
#4	Hikari	thu-coai/seq2seq-pytorch@eff99e	OpenSubtitles (resources://OpenSubtitles_small)	94.698	460f26
#3	Hikari	thu-coai/seq2seq-pytorch@5860c0	OpenSubtitles (resources://OpenSubtitles)	43.868	104528

# Publish Experiments

- Dashboard
  - Compare with others
  - Manage your experiments
  - Hold competitions

# Reproduce Experiments

- Download code from dashboard
  - cotk download ID
- Download code from github
  - cotk download USERNAME/REPO/COMMIT
- Cotk recover codes, cmdlines
  - but not weights
  - Publishers should make their codes reproducible

# Predefined models

- `cotk download thu-coai/seq2seq-pytorch/master`
- `LanguageModel`
- `VAE`
- `SeqGAN`
- `Seq2seq`
- `HRED`
- `CVAE`



Extending CoTK

# New dataset

- Use existing dataloader & Write your preprocessor
  - `LanguageGeneration("./path/to/your_data#your_processor")`
- Override the base dataloader
  - `Class MyDataloader(LanguageGeneration)`
  - Just define some fields, don't need to build vocab list again

# New metrics

- Override MetricBase
  - Implement `__init__`, `forward`, `close`

```
class AverageLengthMetric(MetricBase):
    def __init__(self, dataloader, gen_key="gen"):
        super().__init__()
        self.dataloader = dataloader
        self.gen_key = gen_key
        self.token_num = 0
        self.sent_num = 0

    def forward(self, data):
        gen = data[gen_key]
        for sent in gen:
            self.token_num += len(self.dataloader.trim_index(sent))
            self.sent_num += 1

    def close(self):
        return {"len_avg": self.token_num / self.sent_num}
```

# New models

- Just make your model reproducible
- And upload it to dashboard
- We'll check whether the report is reproducible
- You can become one of our baseline

# Contribution to CoTK

- We will create a package named `cotk_contrib`
- Define your datasets, metrics, models
- Every contribution can be made regardless of quality
- Good contribution will be merged into CoTK

# THANK YOU

- cmd: `pip install cotk`
- Github: [github.com/thu-coai/cotk](https://github.com/thu-coai/cotk)

